

Step Functions Quick Tutorial with Use Case Example

AWS Step Functions allow you to build serverless workflows that coordinate AWS services, including **Lambda** and **DynamoDB**. In this tutorial, we'll cover how Step Functions work, types of states, and how to implement the **use case** of orchestrating a **Lambda API call** using **DynamoDB templates**.

Use Case Breakdown:

1. **DynamoDB** stores API templates (URL, headers, payloads).
 2. **Lambda** dynamically renders API requests using Jinja2 and calls external APIs.
 3. **Step Functions** fetch the template from DynamoDB, call the Lambda, and handle responses.
-

Key Components of Step Functions

1. **States:** These are building blocks in Step Functions. Each state represents a unit of work or decision.
2. **State Types:**
 - **Task:** Executes a Lambda function or another AWS service.
 - **Pass:** Passes input to output without modification.
 - **Choice:** Conditional branching (if/else logic).
 - **Parallel:** Runs multiple branches of states simultaneously.
 - **Wait:** Delays execution for a specific time.
 - **Succeed/Fail:** Marks a successful or failed execution.

Step Function Execution Flow

1. **Start State:** The state where execution begins.
 2. **Transitions:** States are connected via the **Next** property, which defines what happens after one state finishes.
 3. **End State:** Marks the end of the workflow, either success (**Succeed**) or failure (**Fail**).
-

Example: Orchestrating a Lambda API Call Using DynamoDB Template

This example orchestrates the following:

- **Fetch API template** from DynamoDB.
- **Pass template** to Lambda to call an external API.
- **Handle API responses** and flow control.

Step Function Definition

```

{
  "StartAt": "FetchTemplateFromDynamoDB",
  "States": {
    "FetchTemplateFromDynamoDB": {
      "Type": "Task",
      "Resource": "arn:aws:states:::dynamodb:getItem",
      "Parameters": {
        "TableName": "API_Template_Table",
        "Key": {
          "PK": {
            "S.$": "$.templateName"
          },
          "SK": {
            "S.$": "$.method"
          }
        }
      },
      "ResultPath": "$.template",
      "Next": "InvokeLambdaToCallAPI"
    },
    "InvokeLambdaToCallAPI": {
      "Type": "Task",
      "Resource":
"arn:aws:lambda:us-east-1:123456789012:function:GenericApiCallLambda",
      "Parameters": {
        "template.$": "$.template.Item",
        "params.$": "$.params"
      },
      "ResultPath": "$.apiResponse",
      "Next": "ChoiceState"
    },
    "ChoiceState": {
      "Type": "Choice",
      "Choices": [
        {
          "Variable": "$.apiResponse.statusCode",
          "NumericEquals": 200,
          "Next": "Success"
        },
        {

```

```

        "Variable": "$.apiResponse.statusCode",
        "NumericGreaterThanEquals": 400,
        "Next": "HandleFailure"
    }
]
},
"HandleFailure": {
    "Type": "Fail",
    "Cause": "API call failed"
},
"Success": {
    "Type": "Succeed"
}
}
}

```

Explanation of Flow:

1. **FetchTemplateFromDynamoDB** (Task State):
 - Fetches the template from DynamoDB based on a key (**PK** = Template Name, **SK** = HTTP method).
2. **InvokeLambdaToCallAPI** (Task State):
 - Passes the fetched template to a **Lambda** function which dynamically renders and calls the API.
3. **ChoiceState** (Choice State):
 - Based on the API response (**statusCode**), the flow decides:
 - If the response is **200**, it transitions to **Success**.
 - If the response is **>= 400**, it transitions to **HandleFailure**.
4. **Success** (Succeed State):
 - Marks the completion of the flow if the API call was successful.
5. **HandleFailure** (Fail State):
 - Marks the workflow as failed if the API call failed.

Types of States in the Example:

1. **Task State**: Used to perform an action, such as calling Lambda or DynamoDB.
2. **Choice State**: Used to make decisions based on conditions, allowing branching logic.
3. **Succeed State**: Ends the workflow with a success status.
4. **Fail State**: Ends the workflow with a failure status.

Steps to Deploy:

1. **Create the Lambda Function** to dynamically call APIs.
 2. **Create the DynamoDB Table** to store API templates.
 3. **Create the Step Function** using AWS Console or via **CDK**.
-

Summary of Steps in Your Use Case:

1. **Step Function Flow:**
 - Fetches the API template from DynamoDB.
 - Calls a Lambda function to execute the API request.
 - Based on the API response, it either succeeds or fails.
2. **DynamoDB Templates:**
 - Store dynamic API request information (URL, headers, and payload).
3. **Lambda Function:**
 - Uses **Jinja2** to dynamically generate the request and make the API call.

This setup is ideal for coordinating external API calls dynamically with templates stored in DynamoDB and invoking Lambda functions in a serverless architecture.

AWS CDK Core Concepts

1. **Constructs:**
 - Constructs are the basic building blocks of AWS CDK apps. They encapsulate AWS resources and configuration.
 - There are three levels of constructs:
 - **L1 Constructs:** Direct CloudFormation resource mappings.
 - **L2 Constructs:** High-level abstractions of AWS resources (e.g., S3 Bucket, Lambda Function).
 - **L3 Constructs:** Patterns that include multiple L2 resources (e.g., API Gateway + Lambda + DynamoDB).

Example in Python (L2 Construct):

```
from aws_cdk import aws_s3 as s3

bucket = s3.Bucket(self, "MyBucket", versioned=True)
```

2. **App:**

- The CDK App is the root of the CDK application, which contains one or more **stacks**.
- You define an app and initialize stacks inside it.

```
from aws_cdk import core

app = core.App()

stack = core.Stack(app, "MyStack")

app.synth()
```

3. **Stack:**

- A stack is a deployment unit in AWS CDK, representing an AWS CloudFormation stack. Each stack defines the AWS resources that will be deployed together.
- You can have multiple stacks within a single app.

```
from aws_cdk import core

class MyStack(core.Stack):

    def __init__(self, scope: core.Construct, id: str, **kwargs):

        super().__init__(scope, id, **kwargs)

        # Define resources here

app = core.App()

MyStack(app, "MyStack")

app.synth()
```

4. **Resources:**

- AWS resources like S3 buckets, DynamoDB tables, and Lambda functions are defined within a stack.
- These resources are created using **L2 Constructs**.

```

from aws_cdk import aws_s3 as s3, core

class MyStack(core.Stack):

    def __init__(self, scope: core.Construct, id: str, **kwargs):
        super().__init__(scope, id, **kwargs)

        bucket = s3.Bucket(self, "MyBucket", versioned=True)

```

5. Environment:

- CDK allows you to specify the **account** and **region** for deployment. If not specified, it uses the AWS CLI credentials for the current environment.

```

from aws_cdk import core

app = core.App()

MyStack(app, "MyStack", env={'region': 'us-east-1', 'account':
'123456789012'})

app.synth()

```

6. Construct Tree:

- Constructs in CDK form a hierarchical tree. This hierarchy reflects the logical organization of your infrastructure and its scope.

7. Assets:

- Assets represent local files or directories that AWS CDK will upload to S3 and make available for your AWS resources (e.g., Lambda function code).

```

from aws_cdk import aws_lambda as lambda_, core

class MyStack(core.Stack):

    def __init__(self, scope: core.Construct, id: str, **kwargs):
        super().__init__(scope, id, **kwargs)

        lambda_function = lambda_.Function(

```

```
self, "MyLambdaFunction",

runtime=lambda_.Runtime.PYTHON_3_8,

handler="lambda_function.handler",

code=lambda_.Code.from_asset("path/to/lambda/code")

)
```

CDK Features

1. High-Level Abstractions (L2):

- L2 constructs provide higher-level abstractions that simplify resource management.
- For example, creating an **S3 bucket** with versioning enabled can be done with one line.

```
bucket = s3.Bucket(self, "MyBucket", versioned=True)
```

2. Infrastructure as Code:

- CDK abstracts CloudFormation, letting you use familiar programming languages to define infrastructure.
- You can create reusable components and simplify complex configurations.

3. Deployment:

- AWS CDK uses **CloudFormation** under the hood. When you deploy an AWS CDK app, it synthesizes a CloudFormation template and deploys it.
- Use `cdk deploy` to deploy resources.

CDK CLI Commands

cdk init: Initializes a new CDK project.

```
cdk init app --language python
```

cdk synth: Synthesizes the CloudFormation template for your app.

```
cdk synth
```

cdk deploy: Deploys the app to AWS (via CloudFormation).

```
cdk deploy
```

cdk destroy: Destroys the deployed resources.

```
cdk destroy
```

cdk diff: Shows the difference between your local code and the currently deployed stack.

```
cdk diff
```

CDK Use Case Example: Lambda with API Gateway and DynamoDB

This example creates a **Lambda function** connected to **API Gateway** and stores data in **DynamoDB**.

```
from aws_cdk import core

from aws_cdk import aws_lambda as lambda_, aws_apigateway as apigateway,
aws_dynamodb as dynamodb

class MyApiLambdaStack(core.Stack):

    def __init__(self, scope: core.Construct, id: str, **kwargs):
        super().__init__(scope, id, **kwargs)

        # DynamoDB Table

        table = dynamodb.Table(self, "MyTable",
                                partition_key=dynamodb.Attribute(name="ID",
                                                                    type=dynamodb.AttributeType.STRING))

        # Lambda Function

        lambda_function = lambda_.Function(self, "MyLambdaFunction",
                                            runtime=lambda_.Runtime.PYTHON_3_8,
                                            handler="lambda_function.handler",
                                            code=lambda_.Code.from_asset("lambda"))

        # API Gateway

        api = apigateway.LambdaRestApi(self, "MyApi", handler=lambda_function)

        # Grant Lambda access to DynamoDB

        table.grant_read_write_data(lambda_function)

app = core.App()

MyApiLambdaStack(app, "MyApiLambdaStack")

app.synth()
```

Explanation:

1. **Lambda Function:** The function that interacts with DynamoDB.
 2. **API Gateway:** Exposes the Lambda function as a REST API.
 3. **DynamoDB Table:** Stores the data processed by Lambda.
 4. **Permissions:** Lambda is granted permissions to read and write to DynamoDB.
-

Interview Preparation Tips

1. **Understand Core Concepts:**
 - **Constructs, Stacks, Apps, Assets.**
 - CDK lifecycle: **synth, deploy, destroy.**
2. **Hands-on Experience:**
 - Practice creating and deploying CDK apps.
 - Understand how CDK works with services like **Lambda, API Gateway, DynamoDB, S3**, etc.
3. **Common Questions:**
 - How does CDK differ from raw CloudFormation?
 - What are the different levels of constructs in CDK?
 - How does CDK handle environments and regions?
4. **Use CDK in Different Languages:**
 - CDK supports multiple languages (Python, TypeScript, JavaScript, Java).
 - Practice using CDK in your preferred language.

1. Understand Core Concepts

a. Constructs:

- Constructs are the basic building blocks of AWS CDK apps. They represent AWS resources or higher-level abstractions of these resources.
- **L1 Constructs:** Direct mappings to CloudFormation resources, like `CfnBucket` for an S3 bucket. You get full control but need to manage everything yourself.
- **L2 Constructs:** Higher-level, more abstracted constructs (e.g., `s3.Bucket`) that include sensible defaults, making it easier to work with AWS services.
- **L3 Constructs:** Pre-configured patterns that may contain multiple L2 constructs bundled together to solve specific problems. For example, `aws-ecs-patterns.LoadBalancedFargateService` deploys ECS tasks, load balancer, and networking.

b. Stacks:

- A **stack** is a collection of resources defined in your CDK app that are deployed as a unit.
- A stack translates into a CloudFormation stack and can contain multiple AWS resources (Lambda functions, S3 buckets, API Gateway, etc.).
- You can have multiple stacks in a single CDK app. Each stack can be deployed separately.

c. Apps:

- The **App** in CDK is the root of the CDK application and contains the stacks. It is what orchestrates and synthesizes the stacks to deploy resources.
- An app typically consists of one or more stacks.

d. Assets:

- **Assets** represent local files or directories (like Lambda code or container images) that are bundled and uploaded to S3 during deployment.
- The asset can be referenced by AWS services like Lambda or ECS.

2. CDK Lifecycle: synth, deploy, destroy

The lifecycle of working with AWS CDK can be summarized in these three commands:

cdk synth (Synthesize):

- This command generates a CloudFormation template from your CDK code without actually deploying anything. It shows you the CloudFormation template that will be deployed.
- Useful for previewing what your infrastructure will look like in AWS.

`cdk synth`

cdk deploy (Deploy):

- Deploys the resources defined in your CDK app to your AWS account.
- It first generates a CloudFormation template and then uses AWS CloudFormation to deploy the resources.

`cdk deploy`

cdk destroy (Destroy):

- Deletes the resources that were created during `cdk deploy`.
- This command removes all resources that were created as part of the deployment of the specific stack.

`cdk destroy`

How does CDK differ from raw CloudFormation?

- **CDK** provides a higher-level abstraction over CloudFormation. It allows you to write infrastructure as code using familiar programming languages (Python, TypeScript, JavaScript, Java).
- CDK **automatically manages dependencies** between resources and provides high-level **L2 constructs** that come with sensible defaults, making it easier to work with AWS services.
- With **CloudFormation**, you define resources using JSON or YAML, and it's more manual and verbose. CDK simplifies this process with reusable code and constructs.

What are the different levels of constructs in CDK?

- **L1 Constructs**: Direct representations of CloudFormation resources. You need to configure most properties manually (e.g., `s3.CfnBucket`).
- **L2 Constructs**: More abstracted resources with built-in defaults and logic (e.g., `s3.Bucket`). These constructs manage common tasks like setting permissions or managing resources.
- **L3 Constructs**: Pre-configured patterns that implement common solutions (e.g., load-balanced ECS services, serverless APIs).

How does CDK handle environments and regions?

- **Environment** refers to the AWS account and region where the resources will be deployed.
-

You can specify the environment explicitly in your stack:

```
MyStack(app, "MyStack", env={'region': 'us-west-2', 'account':  
'123456789012'})
```

If not specified, CDK will use the credentials and region from the AWS CLI or environment variables.

Core Concepts of AWS Lambda

1. **What is AWS Lambda?**
 - **AWS Lambda** is a serverless compute service that automatically manages the underlying infrastructure for running code.
 - You write the function, deploy it to Lambda, and Lambda automatically scales the function in response to the incoming request load.
2. **Key Features of Lambda:**

- **Event-driven:** Lambda functions are triggered by events from other AWS services (e.g., S3, DynamoDB, API Gateway, CloudWatch).
 - **Stateless:** Lambda is designed to be stateless, meaning each execution is independent. Persistent state must be stored externally (e.g., DynamoDB, S3).
 - **Short-lived executions:** Maximum execution time for a Lambda function is 15 minutes.
3. **Lambda Event Sources:**
- **API Gateway:** Use Lambda to process RESTful requests or expose functions as API endpoints.
 - **S3:** Trigger Lambda functions in response to object uploads, deletions, or changes in a bucket.
 - **DynamoDB:** Respond to changes in DynamoDB tables (e.g., insert, update, delete) using streams.
 - **CloudWatch Events:** Automate tasks based on events like scheduled times or log monitoring.
4. **Lambda Triggers:**
- **Push-based:** Lambda is triggered directly by a service (e.g., API Gateway, S3).
 - **Poll-based:** Lambda polls the event source (e.g., DynamoDB Streams, SQS) and processes the records in the background.
5. **Function Configuration:**
- **Memory:** Lambda functions are configured with memory allocation (128 MB to 10 GB). The CPU scales proportionally with memory.
 - **Timeout:** Set the maximum execution time (up to 15 minutes).
 - **Environment Variables:** Pass configuration details into Lambda functions without hard-coding them.
 - **IAM Role:** Attach an **IAM role** to the Lambda function, allowing it to securely access AWS resources (e.g., S3, DynamoDB).
6. **Concurrency and Scaling:**
- **Concurrency:** Lambda scales automatically with incoming requests. Each request runs in its own environment.
 - **Reserved Concurrency:** You can limit the number of concurrent executions for a Lambda function to prevent overloading downstream resources.
 - **Provisioned Concurrency:** Ensures that a certain number of Lambda instances are initialized and ready to serve requests, improving startup time for cold starts.
7. **Cold Start vs. Warm Start:**
- **Cold Start:** When a Lambda function is invoked for the first time or after a period of inactivity, it experiences a cold start, which involves initialization (slower response).
 - **Warm Start:** Subsequent invocations use already-initialized instances, resulting in faster execution.
8. **Lambda Layers:**
- Layers allow you to share common code or dependencies (e.g., libraries, binaries) across multiple Lambda functions without packaging them in each deployment.
 - Example: Shared utility code or common SDKs can be included as layers.

9. Monitoring and Logging:

- **CloudWatch Logs:** Lambda automatically integrates with CloudWatch Logs for logging execution details (e.g., request ID, start/end time).
- **CloudWatch Metrics:** Monitor execution times, error counts, and request rates for your Lambda functions.

10. Error Handling and Retries:

- **Synchronous Invocations** (e.g., API Gateway): The error is returned to the client. You can handle errors with retries or custom error handling.
 - **Asynchronous Invocations** (e.g., S3, SNS): Failed executions are retried twice, with a delay between attempts.
 - **Dead Letter Queue (DLQ):** Configure DLQs (e.g., SQS, SNS) for failed executions to capture events that failed even after retries.
-

Good to Know Concepts for Lambda

1. Security Best Practices:

- **Least privilege IAM roles:** Ensure Lambda functions have the least amount of access needed to perform their operations.
- **VPC Access:** Lambda can be configured to run inside a VPC to access private resources (e.g., RDS) while still using managed networking features for outbound access.

2. Optimizing Cold Starts:

- Use **Provisioned Concurrency** to keep Lambda instances warm for critical applications.
- Reduce the size of deployment packages to minimize cold start times.

3. Step Functions:

- Lambda functions are often used with AWS Step Functions to build workflows that coordinate multiple functions and services.

4. API Gateway with Lambda:

- Lambda is frequently used in combination with **API Gateway** to create RESTful APIs that dynamically invoke Lambda functions in response to HTTP requests.

5. Versioning and Aliases:

- Lambda supports versioning, allowing you to deploy different versions of a function. **Aliases** can be used to point to specific versions (e.g., **prod**, **dev**).
-

Common Lambda Interview Questions

1. What are cold starts in Lambda, and how do you mitigate them?

- Cold starts occur when a new instance of a Lambda function is initialized. You can mitigate them using **Provisioned Concurrency**, smaller package sizes, and reducing dependency loading.

2. How does Lambda handle scaling?

- Lambda scales automatically by creating multiple instances of the function in response to incoming requests. Each request is processed in parallel, with no need for manual scaling.
- 3. **What are Lambda Layers, and why would you use them?**
 - Layers are shared code or dependencies (libraries) that can be used across multiple Lambda functions. They help avoid duplication and reduce the size of deployment packages.
- 4. **How do you integrate Lambda with other AWS services?**
 - Lambda can be triggered by many AWS services like **S3** (file uploads), **API Gateway** (HTTP requests), **SNS** (notifications), **DynamoDB Streams**, etc.
- 5. **What happens when a Lambda function fails?**
 - For synchronous invocations (e.g., API Gateway), an error response is returned immediately. For asynchronous invocations (e.g., S3), the function is retried twice, and failed requests can be routed to a **Dead Letter Queue (DLQ)**.

1. Overview of DynamoDB

DynamoDB is a fully managed NoSQL database service that provides fast, predictable performance with seamless scalability. It's designed for **high-scale web applications** that require low-latency access to data.

2. Key Concepts in DynamoDB

Tables

- A **DynamoDB table** is where data is stored. Each item in a table is uniquely identified by a primary key.
- Unlike traditional relational databases, there are no columns defined for DynamoDB tables. Each item can have different attributes.

Primary Key

- Every item in a DynamoDB table must have a **primary key**. There are two types:
 - **Partition Key (PK)**: A single attribute that uniquely identifies an item.
 - **Composite Key**: Combines a **Partition Key (PK)** and a **Sort Key (SK)**. The PK identifies the partition, and the SK determines the order of items within the partition.

- **Partition Key Only:** Suitable when items are unique by a single key (e.g., `UserID`).
- **Partition Key and Sort Key:** Used when you want to store multiple items with the same partition key but need an additional level of differentiation (e.g., `UserID` as PK and `OrderID` as SK).

Example

PK (UserID)	SK (OrderID)	Amount	Status
User1	Order1	\$200	Delivered
User1	Order2	\$100	Pending
User2	Order1	\$300	Shipped

In this table, you can retrieve all orders for a specific user (`UserID = User1`), and you can query for a specific order (`UserID = User1, OrderID = Order1`).

3. DynamoDB Keys and Indexes

Partition Key (PK)

- Determines the partition (or physical location) where the item is stored.
- A **partition key** should be high cardinality (unique) to evenly distribute the data.

Sort Key (SK)

- A **sort key** allows you to store multiple items with the same partition key and query them in a sorted order.
- Useful for creating **time-ordered events**, **versioning**, or **grouping related data**.

Global Secondary Index (GSI)

- GSIs allow you to create additional query patterns. You can define a new set of partition and sort keys for querying your table without affecting the primary key.

Example:

- Primary key: `UserID` (PK), `OrderID` (SK)
- GSI: `Status` (PK), `Amount` (SK)
- With the GSI, you can query by `Status` to get all `Pending` orders.

Local Secondary Index (LSI)

- LSIs allow you to add an additional sort key to the primary key (i.e., the partition key stays the same).
- Example: You can query by `UserID` and sort by `OrderDate`.

4. Access Patterns

When designing DynamoDB tables, you need to focus on **access patterns** rather than data normalization. DynamoDB works best when you design your table to fit specific query use cases.

Common Access Patterns

Single Item Query by Partition Key:

python

Copy code

```
response = table.get_item(Key={'UserID': 'User1', 'OrderID':  
'Order1'})
```

-

Query by Partition Key and Sort Key Range:

python

Copy code

```
response = table.query(  
    KeyConditionExpression=Key('UserID').eq('User1') &  
    Key('OrderID').between('Order1', 'Order5')  
)
```

-

Scan (Not Recommended for Large Datasets):

python

Copy code

```
response = table.scan(FilterExpression=Attr('Status').eq('Pending'))
```

-

Global Secondary Index Query:

python

Copy code

```
response = table.query(  
    IndexName='StatusIndex',  
    KeyConditionExpression=Key('Status').eq('Pending')  
)
```

-

5. Best Practices for Web Development

Design for Querying, Not Normalization

- **DynamoDB** is not a relational database. Denormalization and redundant data storage are common and encouraged to support efficient queries.
- Example: If you need fast lookups for both users and orders, store user details in both the **User** table and the **Order** table.

Use Composite Keys for Efficient Queries

- Use **partition keys** and **sort keys** to model hierarchical data (e.g., User -> Orders) and enable efficient querying by ranges.

Leverage GSIs for Additional Query Patterns

- GSIs allow for flexibility in querying the same data in different ways, which is essential when your application needs to support multiple types of queries.

Avoid Scans in Production

- Scans read every item in the table and can be slow and costly. Always prefer **query** operations with defined keys.

Provisioned vs. On-Demand Capacity

- For production apps with predictable workloads, **Provisioned Capacity** allows you to set read/write limits. For unpredictable workloads, **On-Demand Capacity** ensures the table scales automatically without managing capacity.

DynamoDB Streams for Real-Time Processing

- DynamoDB Streams capture table changes (e.g., inserts, updates, deletes) and can trigger AWS Lambda functions for real-time processing, allowing you to react to changes in your application.

Time-Series Data

- For time-series data (e.g., logs, events), use a partition key combined with a timestamp in the sort key to store and query data efficiently.

TTL (Time to Live)

- Use TTL to automatically expire items and remove them from your DynamoDB table after a specified time, useful for session management or expiring cached data.

6. Example Use Case: E-Commerce Orders

Table Schema:

- **PK:** `UserID` (Partition Key)
- **SK:** `OrderID` (Sort Key)
- **Attributes:** `OrderDate`, `Amount`, `Status`

Access Patterns:

Get all orders for a specific user:

python

Copy code

```
response = table.query(
    KeyConditionExpression=Key('UserID').eq('User123')
)
```

1.

Get a specific order for a user:

python

Copy code

```
response = table.get_item(Key={'UserID': 'User123', 'OrderID':
'Order567'})
```

2.

Query orders by status (using GSI):

python

Copy code

```
response = table.query(
    IndexName='StatusIndex',
    KeyConditionExpression=Key('Status').eq('Pending')
)
```

3.

7. Important DynamoDB Features for Web Development

- **Event-Driven Architecture:** Use DynamoDB Streams to trigger real-time updates (e.g., notify users when their order status changes).
 - **Caching with DAX (DynamoDB Accelerator):** Use **DAX** to cache frequently accessed data and reduce latency.
 - **Security:** Use **IAM roles** and policies to secure access to your DynamoDB tables. Enable **encryption at rest** for sensitive data.
-

Summary

For an experienced web developer, **DynamoDB** is a powerful NoSQL database for building scalable, fast web applications. Focus on:

- Designing data models with **PK/SK** to match your query patterns.
- Using **GSIs** to support multiple query patterns efficiently.
- Avoiding expensive operations like **scan**.
- Leveraging **streams** and **TTL** for real-time processing and automatic cleanup.

DynamoDB's flexibility and performance make it ideal for applications that require high availability and low-latency data access.

Let me know if you'd like more details or specific examples for your use case!